



# Parallel Programming and Debugging with CUDA C

**Geoff Gerfin**  
**Sr. System Software Engineer**

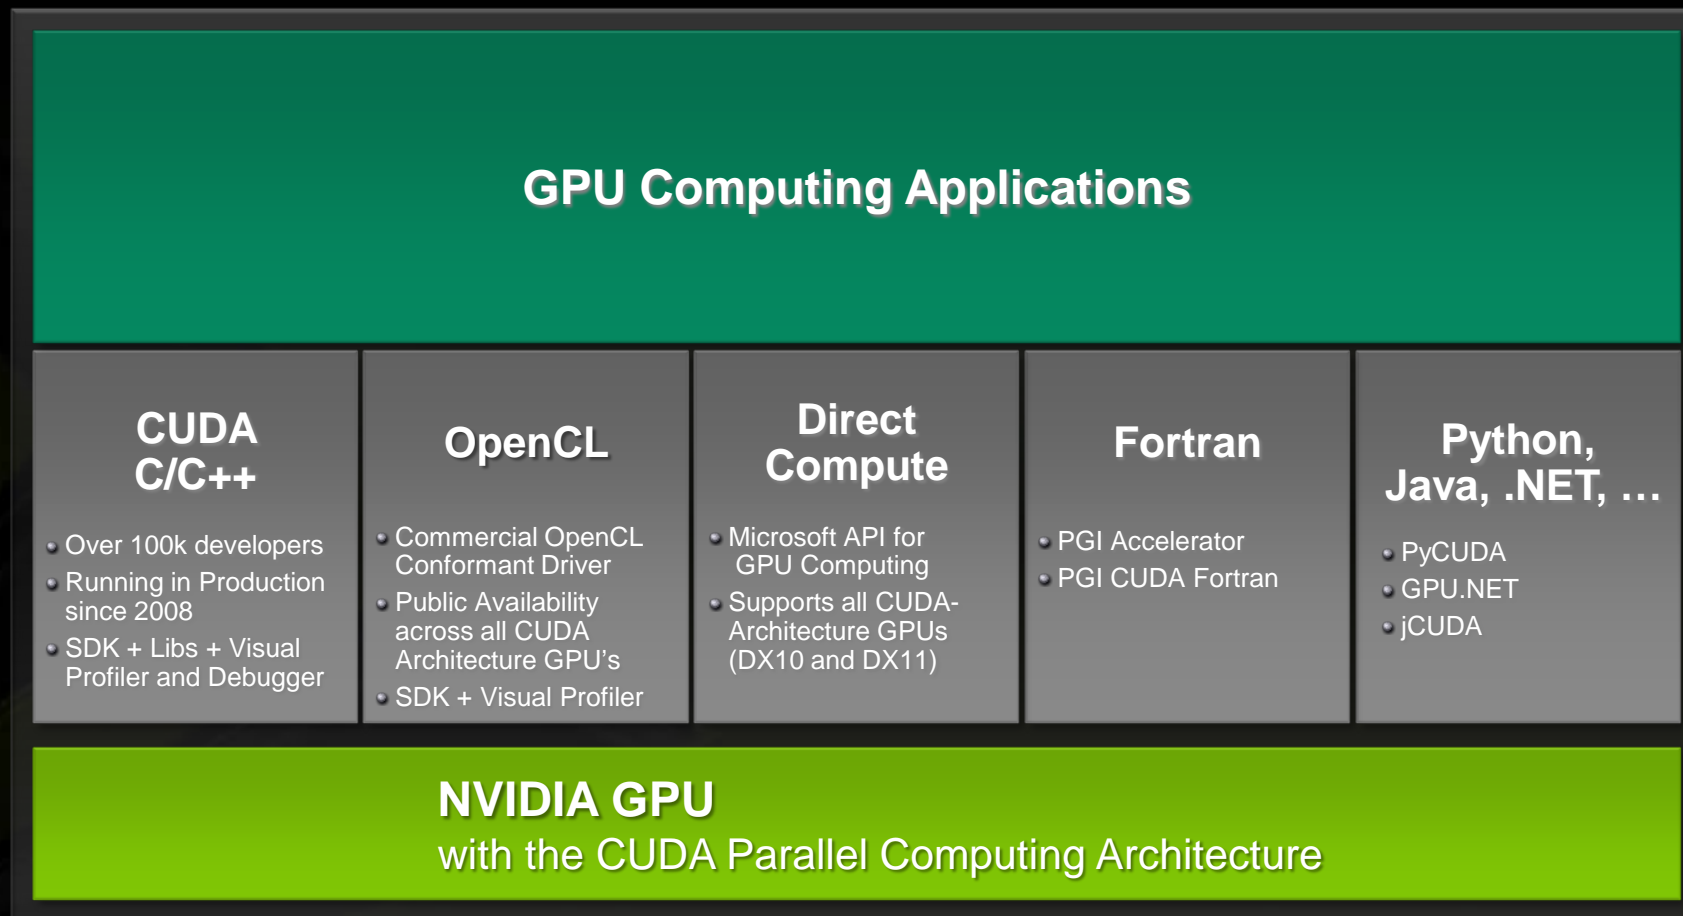


# CUDA - NVIDIA's Architecture for GPU Computing



## Broad Adoption

- Over 250M installed CUDA-enabled GPUs
- Over 650k CUDA Toolkit downloads in last 2 Yrs
- Windows, Linux and MacOS Platforms supported
- GPU Computing spans HPC to Consumer
- 350+ Universities teaching GPU Computing on the CUDA Architecture



- **What will you learn today?**
  - **Write and launch CUDA C kernels**
  - **Manage GPU memory**
  - **Run parallel kernels in CUDA C**
  - **Parallel communication and synchronization**
  - **Debug with cuda-gdb, the Linux CUDA debugger**


# CUDA C: The Basics



- Terminology

- **Host** – The CPU and its memory (host memory)
- **Device** – The GPU and its memory (device memory)

Host



Device



*Note: figure not to scale*



# Hello, World!



```
int main( void ) {  
    printf( "Hello, World!\n" );  
    return 0;  
}
```

- This basic program is just standard C that runs on the *host*
- NVIDIA's compiler, `nvcc`, will not complain about CUDA programs with no *device code*
- At its simplest, CUDA C is just C!

# A Simple Example

- A simple kernel to add two integers:

```
__global__ void add( int *a, int *b, int *c ) {  
    *c = *a + *b;  
}
```

- **CUDA C keyword `__global__` indicates that the `add()` function**
  - Runs on the *device*
  - Called from *host* code

# A Simple Example



- Notice that we use pointers for our variables

```
__global__ void add( int *a, int *b, int *c) {  
    *c = *a + *b;  
}
```

# A Simple Example

- Notice that we use pointers for our variables

```
__global__ void add( int *a, int *b, int *c) {  
    *c = *a + *b;  
}
```

- `add( )` runs on the device...so `a`, `b`, and `c` must point to device memory
- How do we allocate device memory?



# Memory Management



- **Host and device memory are distinct entities**
  - **Device pointers point to GPU memory**
    - May be passed to and from host code
    - May not be dereferenced from host code
  - **Host pointers point to CPU memory**
    - May be passed to and from device code
    - May not be dereferenced from device code



# Memory Management



- **Host and device memory are distinct entities**
  - **Device pointers point to GPU memory**
    - May be passed to and from host code
    - May not be dereferenced from host code
  - **Host pointers point to CPU memory**
    - May be passed to and from device code
    - May not be dereferenced from device code
- **Basic CUDA API for dealing with device memory**
  - `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`
  - **Similar to their C equivalents**, `malloc()`, `free()`, `memcpy()`



# A Simple Example: `add( )`

- **Using our `add( )` kernel:**

```
__global__ void add( int *a, int *b, int *c ) {  
    *c = *a + *b;  
}
```

- **Let's take a look at `main( )...`**

# A Simple Example: `main()`



```
int main( void ) {  
    int a, b, c;           // host copies of a, b, c  
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c  
    int size = sizeof( int ); // we need space for an integer  
}
```

# A Simple Example: `main()`



```
int main( void ) {  
    int a, b, c;           // host copies of a, b, c  
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c  
    int size = sizeof( int ); // we need space for an integer  
  
    // allocate device copies of a, b, c  
    cudaMalloc( (void**)&dev_a, size );  
    cudaMalloc( (void**)&dev_b, size );  
    cudaMalloc( (void**)&dev_c, size );
```

# A Simple Example: `main()`



```
int main( void ) {  
    int a, b, c;           // host copies of a, b, c  
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c  
    int size = sizeof( int ); // we need space for an integer  
  
    // allocate device copies of a, b, c  
    cudaMalloc( (void**)&dev_a, size );  
    cudaMalloc( (void**)&dev_b, size );  
    cudaMalloc( (void**)&dev_c, size );  
  
    a = 2;  
    b = 7;
```



# A Simple Example: `main()` (cont.)



```
// copy inputs to device  
cudaMemcpy( dev_a, &a, size, cudaMemcpyHostToDevice );  
cudaMemcpy( dev_b, &b, size, cudaMemcpyHostToDevice );  
  
}
```

# A Simple Example: `main()` (cont.)



```
// copy inputs to device  
cudaMemcpy( dev_a, &a, size, cudaMemcpyHostToDevice );  
cudaMemcpy( dev_b, &b, size, cudaMemcpyHostToDevice );  
  
// launch add() kernel on GPU, passing parameters  
add<<< 1, 1 >>>( dev_a, dev_b, dev_c );
```

# A Simple Example: `main()` (cont.)



```
// copy inputs to device
cudaMemcpy( dev_a, &a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, &b, size, cudaMemcpyHostToDevice );

// launch add() kernel on GPU, passing parameters
add<<< 1, 1 >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( &c, dev_c, size, cudaMemcpyDeviceToHost );

}
```

# A Simple Example: `main()` (cont.)



```
// copy inputs to device
cudaMemcpy( dev_a, &a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, &b, size, cudaMemcpyHostToDevice );

// launch add() kernel on GPU, passing parameters
add<<< 1, 1 >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( &c, dev_c, size, cudaMemcpyDeviceToHost );

cudaFree( dev_a );
cudaFree( dev_b );
cudaFree( dev_c );
```

# A Simple Example: `main()` (cont.)



```
// copy inputs to device
cudaMemcpy( dev_a, &a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, &b, size, cudaMemcpyHostToDevice );

// launch add() kernel on GPU, passing parameters
add<<< 1, 1 >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( &c, dev_c, size, cudaMemcpyDeviceToHost );

cudaFree( dev_a );
cudaFree( dev_b );
cudaFree( dev_c );
return 0;
}
```

# Parallel Programming in CUDA C



- But wait...GPU computing is about massive parallelism
- So how do we run code in parallel on the device?



# Parallel Programming in CUDA C



- But wait...GPU computing is about massive parallelism
- So how do we run code in parallel on the device?
- Solution lies in the parameters between the triple angle brackets:

# Parallel Programming in CUDA C



- But wait...GPU computing is about massive parallelism
- So how do we run code in parallel on the device?
- Solution lies in the parameters between the triple angle brackets:

```
add<<< 1, 1 >>>( dev_a, dev_b, dev_c );
```



```
add<<< N, 1 >>>( dev_a, dev_b, dev_c );
```

- Instead of executing `add( )` once, `add( )` executed N times in parallel

# Parallel Programming in CUDA C



- With `add( )` running in parallel, let's do vector addition
- Terminology: Each parallel invocation of `add( )` referred to as a **block**

# Parallel Programming in CUDA C

- With `add( )` running in parallel, let's do vector addition
- Terminology: Each parallel invocation of `add( )` referred to as a **block**
- Kernel can refer to its block's index with variable `blockIdx.x`
- Each block adds a value from `a[ ]` and `b[ ]` , storing the result in `c[ ]`:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

- By using `blockIdx.x` to index arrays, each block handles different indices

# Parallel Programming in CUDA C



- We write this code:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

# Parallel Programming in CUDA C



- We write this code:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

- This is what runs in parallel on the device:

Block 0

```
c[0] = a[0] + b[0];
```

Block 1

```
c[1] = a[1] + b[1];
```

Block 2

```
c[2] = a[2] + b[2];
```

Block 3

```
c[3] = a[3] + b[3];
```



# Parallel Addition: `add( )`



- **Using our newly parallelized `add( )` kernel:**

```
__global__ void add( int *a, int *b, int *c ) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

- **Let's take a look at `main( )...`**

# Parallel Addition: `main()`



```
#define N 512
int main( void ) {
    int *a, *b, *c;           // host copies of a, b, c
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c
```

# Parallel Addition: main( )



```
#define N 512
int main( void ) {
    int *a, *b, *c;           // host copies of a, b, c
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c
    int size = N * sizeof( int ); // we need space for 512 integers

    // allocate device copies of a, b, c
    cudaMalloc( (void**)&dev_a, size );
    cudaMalloc( (void**)&dev_b, size );
    cudaMalloc( (void**)&dev_c, size );

    a = (int*)malloc( size );
    b = (int*)malloc( size );
    c = (int*)malloc( size );

    random_ints( a, N );
    random_ints( b, N );
```

# Parallel Addition: `main( )` (cont.)



```
// copy inputs to device  
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );  
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );
```

# Parallel Addition: `main()` (cont.)



```
// copy inputs to device  
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );  
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );  
  
// launch add() kernel with N parallel blocks  
add<<< N, 1 >>>( dev_a, dev_b, dev_c );
```

# Parallel Addition: `main()` (cont.)



```
// copy inputs to device
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );

// launch add() kernel with N parallel blocks
add<<< N, 1 >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( c, dev_c, size, cudaMemcpyDeviceToHost );
```



# Parallel Addition: `main()` (cont.)



```
// copy inputs to device
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );

// launch add() kernel with N parallel blocks
add<<< N, 1 >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( c, dev_c, size, cudaMemcpyDeviceToHost );

free( a ); free( b ); free( c );
cudaFree( dev_a );
cudaFree( dev_b );
cudaFree( dev_c );
return 0;
```

```
}
```

# Threads



- Terminology: A block can be split into parallel *threads*
- Let's change vector addition to use parallel threads instead of parallel blocks:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[ blockIdx.x ] = a[ blockIdx.x ] + b[ blockIdx.x ];  
}
```

# Threads



- Terminology: A block can be split into parallel **threads**
- Let's change vector addition to use parallel threads instead of parallel blocks:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[ threadIdx.x ] = a[ threadIdx.x ] + b[ threadIdx.x ];  
}
```

- We use `threadIdx.x` instead of `blockIdx.x` in `add()`

# Threads



- Terminology: A block can be split into parallel **threads**
- Let's change vector addition to use parallel threads instead of parallel blocks:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[ threadIdx.x ] = a[ threadIdx.x ] + b[ threadIdx.x ];  
}
```

- We use `threadIdx.x` instead of `blockIdx.x` in `add()`
- `main()` will require one change as well...

# Parallel Addition (Threads): `main()`



```
#define N 512
int main( void ) {
    int *a, *b, *c;           // host copies of a, b, c
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c
    int size = N * sizeof( int ); // we need space for 512 integers

    // allocate device copies of a, b, c
    cudaMalloc( (void**)&dev_a, size );
    cudaMalloc( (void**)&dev_b, size );
    cudaMalloc( (void**)&dev_c, size );

    a = (int*)malloc( size );
    b = (int*)malloc( size );
    c = (int*)malloc( size );

    random_ints( a, N );
    random_ints( b, N );
```

# Parallel Addition (Threads): `main()` (cont.)



```
// copy inputs to device  
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );  
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );
```

# Parallel Addition (Threads): `main()` (cont.)



```
// copy inputs to device
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );

// launch add() kernel with N parallel threads
add<<< 1, N >>>( dev_a, dev_b, dev_c );
```

# Parallel Addition (Threads): `main()` (cont.)



```
// copy inputs to device
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );

// launch add() kernel with N parallel threads
add<<< 1, N >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( c, dev_c, size, cudaMemcpyDeviceToHost );

free( a ); free( b ); free( c );
cudaFree( dev_a );
cudaFree( dev_b );
cudaFree( dev_c );
return 0;
}
```



# Why Bother With Threads?



- **Threads seem unnecessary**
  - Added a level of abstraction and complexity
  - What did we gain?


# Why Bother With Threads?



- **Threads seem unnecessary**
  - Added a level of abstraction and complexity
  - What did we gain?
- **Unlike parallel blocks, parallel threads have mechanisms to:**
  - Communicate
  - Synchronize
- **Let's see how...**

# Dot Product


- Unlike vector addition, dot product is a *reduction* from vectors to a scalar



# Dot Product



- Unlike vector addition, dot product is a *reduction* from vectors to a scalar




$$c = \vec{a} \cdot \vec{b}$$

$$c = (a_0, a_1, a_2, a_3) \cdot (b_0, b_1, b_2, b_3)$$

$$c = a_0 b_0 + a_1 b_1 + a_2 b_2 + a_3 b_3$$

# Dot Product


- Parallel threads have no problem computing the pairwise products:



# Dot Product



- Parallel threads have no problem computing the pairwise products:




- So we can start a dot product CUDA kernel by doing just that:

```
__global__ void dot( int *a, int *b, int *c )    {  
    // Each thread computes a pairwise product  
    int temp = a[threadIdx.x] * b[threadIdx.x];
```

# Dot Product


- But we need to share data between threads to compute the final sum:



# Dot Product



- But we need to share data between threads to compute the final sum:



```
__global__ void dot( int *a, int *b, int *c ) {  
    // Each thread computes a pairwise product  
    int temp = a[threadIdx.x] * b[threadIdx.x];  
  
    // Can't compute the final sum  
    // Each thread's copy of 'temp' is private  
}
```



# Sharing Data Between Threads



- Terminology: A block of threads shares memory called...


# Sharing Data Between Threads



- Terminology: A block of threads shares memory called...*shared memory*

# Sharing Data Between Threads

- Terminology: A block of threads shares memory called...*shared memory*
- Extremely fast, on-chip memory (user-managed cache)
- Declared with the `__shared__` CUDA keyword
- Not visible to threads in other blocks running in parallel



# Parallel Dot Product: dot()



- We perform parallel multiplication, serial addition:

```
#define N 512
__global__ void dot( int *a, int *b, int *c ) {
    // Shared memory for results of multiplication
    __shared__ int temp[N];
    temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
```

# Parallel Dot Product: dot ( )



- We perform parallel multiplication, serial addition:

```
#define N 512
__global__ void dot( int *a, int *b, int *c ) {
    // Shared memory for results of multiplication
    __shared__ int temp[N];
    temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];

    // Thread 0 sums the pairwise products
    if ( 0 == threadIdx.x ) {
        int sum = 0;
        for (int i = 0; i < N; i++)
            sum += temp[i];
        *c = sum;
    }
}
```

# Parallel Dot Product Recap



- We perform parallel, pairwise multiplications

# Parallel Dot Product Recap



- We perform parallel, pairwise multiplications
- Shared memory stores each thread's result

# Parallel Dot Product Recap



- We perform parallel, pairwise multiplications
- Shared memory stores each thread's result
- We sum these pairwise products from a single thread
- Sounds good...



# Parallel Dot Product Recap



- We perform parallel, pairwise multiplications
- Shared memory stores each thread's result
- We sum these pairwise products from a single thread
- Sounds good... but we've made a huge mistake

# Enter the Debugger



- We will demonstrate how **cuda-gdb** can be used to find a bug in our `dot()` kernel

# Enter the Debugger



- We will demonstrate how **cuda-gdb** can be used to find a bug in our `dot()` kernel
- The debugger follows CUDA language semantics when advancing program execution:

# Enter the Debugger



- We will demonstrate how **cuda-gdb** can be used to find a bug in our `dot()` kernel
- The debugger follows CUDA language semantics when advancing program execution:
  - When single-stepping a CUDA thread, the entire **warp** it belongs to will single-step
  - A warp is a group of 32 CUDA threads

# Enter the Debugger



- We will demonstrate how **cuda-gdb** can be used to find a bug in our `dot( )` kernel
- The debugger follows CUDA language semantics when advancing program execution:
  - When single-stepping a CUDA thread, the entire **warp** it belongs to will single-step
  - A warp is a group of 32 CUDA threads
- **Simply tracking how the program advances can reveal synchronization issues**

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c )    {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

**(cuda-gdb)**

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c )    {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

**(cuda-gdb) break dot**

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c ) {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

(cuda-gdb) run



# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c ) {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

**(cuda-gdb) info cuda threads**

**<<<(0,0),(0,0,0)>>> ... <<<(0,0),(511,0,0)>>> at dotproduct.cu:5**

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c )    {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

(cuda-gdb) next

# Debugging with cuda-gdb



```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c )    {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

(cuda-gdb) next

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c )    {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

(cuda-gdb) next

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c )    {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

(cuda-gdb) next

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c )    {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

(cuda-gdb) next

<<<(0,0),(0,0,0)>>> ... <<<(0,0),(0,0,0)>>> at dotproduct.cu:11

<<<(0,0),(1,0,0)>>> ... <<<(0,0),(31,0,0)>>> at dotproduct.cu:14

<<<(0,0),(32,0,0)>>> ... <<<(0,0),(511,0,0)>>> at dotproduct.cu:5

# Debugging with cuda-gdb

```
1  #define N 512
2  __global__ void dot( int *a, int *b, int *c ) {
3      // Shared memory for results of multiplication
4      __shared__ int temp[N];
5      temp[threadIdx.x] = a[threadIdx.x] * b[threadIdx.x];
6      __syncthreads();
7      // Thread 0 sums the pairwise products
8      if ( 0 == threadIdx.x ) {
9          int sum = 0;
10         for (int i = 0; i < N; i++)
11             sum += temp[i];
12         *c = sum;
13     }
14 }
```

Threads 32 through 511 did not write out their results yet. To fix this bug, we need to synchronize all threads in this block.

(cuda-gdb) next

```
<<<(0,0),(0,0,0)>>> ... <<<(0,0),(0,0,0)>>> at dotproduct.cu:11
<<<(0,0),(1,0,0)>>> ... <<<(0,0),(31,0,0)>>> at dotproduct.cu:14
<<<(0,0),(32,0,0)>>> ... <<<(0,0),(511,0,0)>>> at dotproduct.cu:5
```

# NVIDIA cuda-gdb



## CUDA debugging **integrated** into GDB on Linux

- Supported on **32bit** and **64bit** systems
- Seamlessly** debug both the host/CPU and device/GPU code
- Set **breakpoints** on any source line or symbol name
- Access and print all CUDA memory allocs, local, global, constant and shared vars

Included in the CUDA Toolkit

The screenshot displays the NVIDIA cuda-gdb interface, which is integrated into the Emacs editor. The top window shows the host code (C/C++) being debugged, with the current thread at the `acos_main` function. The bottom window shows the device code (CUDA) being debugged, with the current thread at the `__device_func__` function. The interface includes a menu bar (File, Edit, Options, Buffers, Tools, Gdb, Complete, In/Out, Signals, Help) and a toolbar with various debugging symbols. The left pane shows the source code with breakpoints set, and the right pane shows the current thread's state and variables.

```
*gud-acos_dbg* - emacs@ssalian-linux
File Edit Options Buffers Tools Gdb Complete In/Out Signals Help

p p* [debugging icons]

[Current CUDA Thread <<<(0,0),(>
acos_main () at /ssalian-local/
(cuda-gdb) s
[Current CUDA Thread <<<(0,0),(>
acos_main () at /ssalian-local/
Breakpoint 2 at 0x805abc4c: fil>

#else /* FERMI */
acos_main<<<ACOS_CTA_CNT,ACOS_THREAD_CNT>>>
#endif
stop = second();
cudaStat = cudaGetLastError(); /* check f>

NX - ssalian@172.16.175.110:1022 - ssalian-linux
Applications Places System [icons]

*gud-acos_dbg* - emacs@ssalian-linux
File Edit Options Buffers Tools Gdb Complete In/Out Signals Help

p p* [debugging icons]

[Current CUDA Thread <<<(0,0),(>
Breakpoint 1, acos_main () at a>
(cuda-gdb) s
[Current CUDA Thread <<<(0,0),(>
acos_main () at acos.cu:390
(cuda-gdb) s
[Current CUDA Thread <<<(0,0),(>
acos_main () at acos.cu:391
(cuda-gdb) p threadIdx
$5 = {x = 0, y = 0, z = 0}
(cuda-gdb) p blockIdx
$6 = {x = 0, y = 0}
(cuda-gdb) info cuda threads
<<<(0,0),(0,0,0)>>> ... <<<(0,0)
<<<(0,0),(32,0,0)>>> ... <<<(23>
(cuda-gdb) p blockDim
$7 = {x = 128, y = 1, z = 1}

__device_func__(float __cuda_acosf(float a))
{
float t0, t1, t2;

t0 = __cuda_fabsf(a);
t2 = 1.0f - t0;
t2 = 0.5f * t2;
t2 = __cuda_sqrtf(t2);
t1 = t0 > 0.57f ? t2 : t0;
t1 = __internal_asinf_kernel(t1);
t1 = t0 > 0.57f ? 2.0f * t1 : CUDART_PIO2_F>
if (__cuda_signbitf(a)) {
t1 = CUDART_PI_F - t1;
}
#if !defined(__CUDABE__)
if (__cuda_isnanf(a)) {
t1 = a + a;
}
}

$1 = {x = 128, y = 1, z = 1}
(cuda-gdb) b blockDim
<<<(0,0),(32,0,0)>>> ... <<<(32>
<<<(0,0),(0,0,0)>>> ... <<<(0,0)>
```



```
Applications Places System
linux64:~/demo2010$ ./ptrchecktest
unspecified launch failure : 79
linux64:~/demo2010$ cuda-memcheck ./ptrchecktest
===== CUDA-MEMCHECK
unspecified launch failure : 79
===== Invalid __global__ read of size 4
===== at 0x00000158 in ptrchecktest.cu:27:kernel2
===== by thread (0,0,0) in block (0,0)
===== Address 0xfd0000001 is misaligned
=====
===== ERROR SUMMARY: 1 error
linux64:~/demo2010$ cuda-memcheck --continue ./ptrchecktest
===== CUDA-MEMCHECK
Checking...
Done
Checking...
Error: 3 (0)
Done
Checking...
Error: 1 (0)
Error: 3 (0)
Error: 5 (0)
Error: 7 (0)
Done
===== Invalid __global__ read of size 4
===== at 0x00000158 in ptrchecktest.cu:27:kernel2
===== by thread (0,0,0) in block (0,0)
===== Address 0xfd0000001 is misaligned
=====
===== Invalid __global__ read of size 4
===== at 0x00000198 in ptrchecktest.cu:18:kernel1
===== by thread (3,0,0) in block (5,0)
===== Address 0xfd00000028 is out of bounds
=====
===== Invalid __global__ write of size 8
===== at 0x000001d0 in ptrchecktest.cu:38:kernel3
===== by thread (1,0,0) in block (8,0)
===== Address 0xfd00000204 is misaligned
=====
===== Invalid __global__ write of size 4
===== at 0x000000f0 in ptrchecktest.cu:44:kernel4
===== by thread (63,0,0) in block (22,0)
===== Address 0x00000000 is out of bounds
=====
===== ERROR SUMMARY: 4 errors
```

# cuda-memcheck



- **Detect memory and threading errors**
  - OOB memory accesses
  - Misaligned memory accesses
- **Windows, Linux, and Mac OSX**
- **Usage**
  - Standalone: `cuda-memcheck <app>`
  - `cuda-gdb`: set cuda memcheck on

Included in the **CUDA Toolkit**

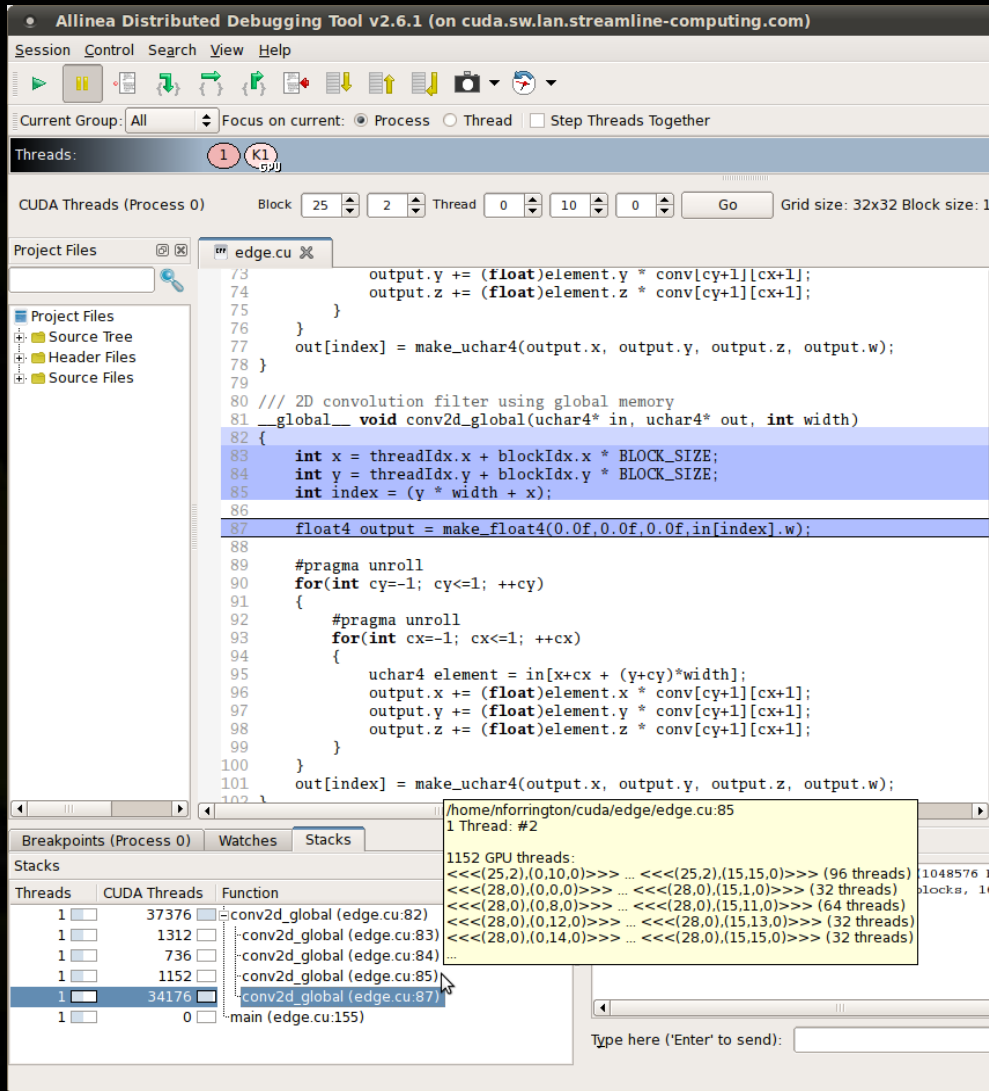
# Parallel Nsight for Visual Studio



- Integrated development for CPU and GPU



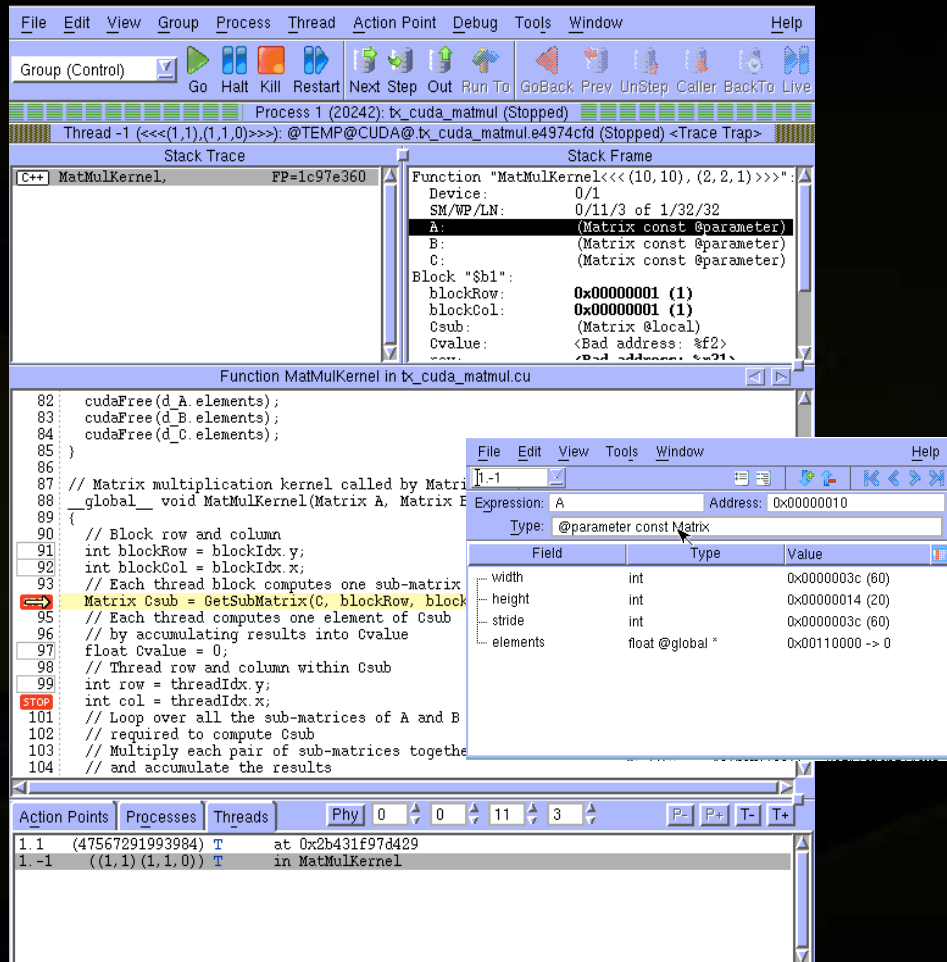
# Allinea DDT Debugger



- Fermi and Tesla support
- cuda-memcheck support for memory errors
- Combined MPI and CUDA support
- Stop on kernel launch feature
- Kernel thread control, evaluation and breakpoints
  - Identify thread counts, ranges and CPU/GPU threads easily
- Multi-Dimensional Array Viewer (MDA)
  - 3D Data Visualization
- Coming soon: multiple GPU device support



# TotalView Debugger



- Full visibility of both Linux threads and GPU device threads
  - Device threads shown as part of the parent Unix process
  - Correctly handle all the differences between the CPU and GPU
- Fully represent the hierarchical memory
  - Display data at any level (registers, local, block, global or host memory)
  - Making it clear where data resides with type qualification
- Thread and Block Coordinates
  - Built in runtime variables display threads in a warp, block and thread dimensions and indexes
  - Displayed on the interface in the status bar, thread tab and stack frame
- Device thread control
  - Warps advance synchronously
- Handles CUDA function inlining
  - Step into or over inlined functions
- Reports memory access errors
  - CUDA memcheck
- Can be used with MPI

# Questions?



- **Latest CUDA Toolkit and Driver**
  - <http://www.nvidia.com/getcuda>
- **Additional Resources on CUDA from GTC 2010**
  - <http://www.nvidia.com/gtc>
- **PGI CUDA C for Multi-Core x86 Processors**
  - Wednesday, 11/17 @ 1:00pm